

Machine Learning: Concepts and Applications






Table of contents

01

AI vs. Machine
Learning vs.
Deep Learning

02

Introduction to
Machine
Learning

03

Types of
Machine
Learning

04

Fundamental
Techniques of
Machine Learning


05

Basic
introduction to
deep learning

06

Introduction to
convolution neural
network (CNN)





01

AI vs. Machine Learning vs. Deep Learning



2 تعلم الآلة (Machine Learning)

التعريف

هو فرع من الذكاء الاصطناعي يجعل النظام يتعلم من البيانات بدلاً من برمجته صراحةً.

كيف يعمل؟

- نعطيه بيانات (مدخلات + مخرجات صحيحة)
- يتعلم الأنماط
- يستخدم نموذجًا رياضيًا للتنبؤ

مثال

تصنيف البريد الإلكتروني (Spam / Not Spam):

- ندخل آلاف الرسائل المصنفة مسبقًا
- النظام يتعلم الكلمات والأنماط
- يقرر تلقائيًا مستقبلًا

هنا النظام يتعلم، لكن ما زال يحتاج استخراج ميزات يدويًا.

- الذكاء الاصطناعي (AI) ← المفهوم الأشمل
- تعلم الآلة (ML) ← جزء من الذكاء الاصطناعي
- التعلم العميق (DL) ← جزء متقدم من تعلم الآلة

1 الذكاء الاصطناعي (Artificial Intelligence)

التعريف

هو المجال الذي يهدف إلى جعل الحاسوب يتصرف بذكاء يشبه الإنسان مثل:

- التفكير
- اتخاذ القرار
- حل المشكلات
- التخطيط والاستدلال

كيف يعمل؟

ليس شرطًا أن يتعلم من البيانات، بل قد يعتمد على:

- قواعد منطقية
- إذا... فإن (IF-THEN)
- محركات استدلال





3 التعلّم العميق (Deep Learning)

التعريف

هو نوع متقدم من تعلّم الآلة يعتمد على الشبكات العصبية العميقة المشابهة لدماغ الإنسان.

كيف يعمل؟

- يتعلّم الميزات بنفسه
- يستخدم طبقات متعددة من الخلايا العصبية
- فعال جدًا مع الصور، الصوت، الفيديو

مثال

التعرّف على الوجوه:

- نعطي النظام صورة فقط
- هو يتعلّم تلقائيًا:
 - الحواف
 - العيون
 - الأنف
 - ملامح الوجه الكاملة

لا نحدد له ما هي العين أو الأنف، بل يكتشفها بنفسه.



AI vs. Machine learning vs. Deep Learning

● Artificial Intelligence


A science devoted to making machines think and act like humans.

● Machine Learning

Focuses on enabling computers to perform tasks without explicit programming.

● Deep Learning

A subset of machine learning based on artificial neural networks.



skiplevel.co



- A good analogy is to imagine AI as transportation encompassing all methods of moving people (e.g. cars, trains, bicycles).
- Machine learning would be like cars, and
- Deep learning would be electric cars.





Artificial Intelligence (AI)



Definition

AI is the broader concept of creating machines or software that can perform tasks typically requiring human intelligence. This includes reasoning, problem-solving, understanding natural language, and perception.



Scope

AI encompasses a wide range of technologies and methodologies, including rule-based systems, expert systems, and various learning algorithms.



Examples

Chatbots, recommendation systems, voice assistants (like Siri and Alexa), and game-playing AI (like AlphaGo).





Machine Learning (ML)



Definition

A subset of AI that focuses on the idea that systems can learn from data, identify patterns, and make decisions with minimal human intervention.



Scope

Machine learning involves algorithms that enable computers to learn from and make predictions based on data. It is not limited to one technique but includes various approaches such as supervised, unsupervised, and reinforcement learning.



Examples

Spam detection in email, image recognition, and predictive analytics (like predicting stock prices).



Deep Learning



Definition

A subset of machine learning that uses neural networks with many layers (hence "deep") to model complex patterns in large amounts of data.



Scope

Deep learning is particularly effective for tasks that involve high-dimensional data such as images, audio, and text. It automatically learns features from raw data, reducing the need for manual feature extraction.



Examples

Image and speech recognition, natural language processing (NLP), and autonomous vehicles.





Summary of Differences


Hierarchy:

- AI is the overarching field.
- Machine learning is a subset of AI.
- Deep learning is a specialized subset of machine learning.

Complexity:

- AI encompasses all intelligent behavior in machines.
- Machine learning focuses on data-driven learning.
- Deep learning tackles highly complex tasks using neural networks.






02

Introduction to Machine Learning



Machine learning workflow



Key Concepts in Machine Learning



Data

- The foundation of machine learning. Data can be structured (like spreadsheets) or unstructured (like images and text).
- Quality and quantity of data significantly influence the performance of ML models.



Features

- Attributes or variables in the data used for making predictions. For instance, in a dataset predicting house prices, features could include size, location, and number of bedrooms.

Key Concepts in Machine Learning2



Labels

- The output variable you want to predict.
- In supervised learning, each input feature set is associated with a label.



Training and Testing Sets

- The process of teaching the model using data.
- The dataset is usually split into training and testing sets. The training set is used to train the model, while the testing set evaluates its performance.

Key Concepts in Machine Learning3




Algorithms

- Various algorithms are used depending on the task. Common types include:
 - **Linear Regression:** For predicting continuous values.
 - **Logistic Regression:** For binary classification tasks.
 - **Decision Trees:** For classification and regression tasks.
 - **Support Vector Machines (SVM):** Effective for high-dimensional spaces.
 - **Neural Networks:** For complex patterns, especially in deep learning.

Machine learning algorithms

Algorithm	Type	Accuracy	Training Time	Applications
Linear Regression	Supervised	Moderate	Fast	Prediction, Regression
Decision Trees	Supervised	High	Moderate	Classification, Regression
Neural Networks	Supervised	High	Slow	Image Processing
K-Means	Unsupervised	Moderate	Fast	Data Clustering
SVM	Supervised	High	Moderate	Classification, Anomaly Detection




03

Types of Machine Learning



Supervised Learning in ML



Supervised Learning

- **Definition:** Learning from labeled data.
- **Example Algorithms:** Linear regression, decision trees, support vector machines.

Applications of supervised learning

Image Classification

Example: Identifying objects in images (e.g., classifying photos of cats and dogs).

How it Works: The model learns to recognize patterns in pixel data associated with different **labels**, improving its accuracy with more training images.

Speech Recognition

Example: Converting spoken language into text.

How it Works: The model is trained on audio samples paired with their transcriptions, allowing it to learn how different sounds correspond to written language.

Applications of supervised learning2



Medical Diagnosis

Example: Predicting whether a patient has a specific disease based on medical history and test results.

How it Works: The model is trained on historical patient data, learning to associate specific symptoms and test results with diagnoses.

Sentiment Analysis

Example: Determining whether customer reviews are positive, negative, or neutral.

How it Works: The model is trained on a dataset of labeled reviews, learning to identify sentiment through linguistic patterns and word usage.



Applications of supervised learning3

Spam Detection

Example: Classifying emails as "spam" or "not spam".


How it Works: A supervised model is trained on a dataset of labeled emails, learning features that typically indicate spam (like certain keywords or formats).

House Price Prediction

Example: Estimating the price of a house based on its features (size, location, number of bedrooms).

How it Works: The model learns from historical sales data, associating various features with sale prices to make future predictions.

Unsupervised Machine Learning



Unsupervised Learning

- **Definition:** Learning from unlabeled data to identify patterns.
- **Example Algorithms:** K-means clustering, hierarchical clustering, principal component analysis (PCA).

Applications of Unsupervised learning

Image Compression

Example: Reducing the size of images while preserving quality.

How it Works: Techniques like clustering can group similar pixels, allowing for more efficient storage by reducing redundancy in image data.

Speech and Voice Recognition

Example: Analyzing speech patterns for various applications.

How it Works: Unsupervised learning can identify features in audio signals without labeled datasets, useful for enhancing voice recognition systems.

Applications of Unsupervised learning2



Customer Segmentation

Example: Grouping customers based on purchasing behavior.

How it Works: Algorithms like K-Means or hierarchical clustering can analyze transaction data to identify distinct customer segments, enabling targeted marketing strategies.

Anomaly Detection

Example: Identifying fraudulent transactions in financial systems.

How it Works: By modeling normal transaction patterns, unsupervised learning can flag unusual transactions that deviate significantly from these patterns, helping to detect fraud.



Applications of Unsupervised learning3

Social Network Analysis


Example: Understanding relationships and influence within social networks.

How it Works: Unsupervised methods can analyze user interactions and connections to identify communities or influential users within a network.

Recommendation Systems

Example: Suggesting products or content based on user behavior.

How it Works: Collaborative filtering can identify patterns in user preferences without explicit labels, recommending items based on similar users' choices.



Reinforcement Learning

- **Definition:** Learning through trial and error, receiving rewards or penalties.
- **Example Algorithms:** Q-learning, deep Q-networks.

Applications of Reinforcement learning



Robotics

Example: Robots learning to navigate and manipulate objects.

How it Works: Robots use reinforcement learning to optimize their movements based on feedback from their environment, improving tasks like walking, grasping, or flying.



Autonomous Vehicles

Example: Self-driving cars learning to navigate through traffic.

How it Works: RL algorithms enable vehicles to learn optimal driving policies by simulating various driving scenarios, receiving rewards for safe driving and penalties for accidents or rule violations.

Applications of Reinforcement learning2



Healthcare

Example: Personalized treatment plans for patients.

How it Works: RL can optimize treatment strategies based on patient responses, learning which interventions yield the best outcomes through trial and error.



Personalized Learning

Example: Adaptive educational platforms.

How it Works: RL can tailor educational content to individual learners, optimizing engagement and knowledge retention by adjusting difficulty and type of material based on learner responses.

Applications of Reinforcement learning3

Natural Language Processing (NLP)

Example: Chatbots and dialogue systems.

How it Works: RL can optimize conversational agents to provide better responses by rewarding successful interactions and penalizing unhelpful ones, enhancing user satisfaction.

Energy Management

Example: Smart grid and energy consumption optimization.

How it Works: RL can manage energy distribution in smart grids by learning optimal consumption patterns and reducing costs while maintaining service levels.

Machine Learning
Techniques with
Python



04


Fundamental Techniques of Machine Learning





Classification

Definition: Classification is a type of supervised learning in machine learning where the goal is to identify which category an object belongs to based on its features.



Key Concepts in Classification



Training Data

- The dataset used to train the model, consisting of features (input variables) and labels (output categories).



Features

- Characteristics or attributes used to describe the data points.

Key Concepts in Classification2



Labels

- Categorical outcomes that the model is trying to predict (e.g., Normal image or Abnormal image).



Model

- The algorithm used to learn from the training data and make predictions.

Common Classification Algorithms




Algorithms

- Linear Regression: For predicting continuous values.
- Logistic Regression: For binary classification tasks.
- Decision Trees: For classification and regression tasks.
- Support Vector Machines (SVM): Effective for high-dimensional spaces.
- Neural Networks: For complex patterns, especially in deep learning.

Medical diagnosis Application



Is a critical application of supervised machine learning, where algorithms are trained on labeled patient data to predict diseases or conditions based on symptoms, test results, and other clinical features.



Key Components of Medical Diagnosis with Supervised Learning



Data Collection

- Electronic Health Records (EHRs): Patient data including demographics, medical history, laboratory test results, and treatment outcomes.
- Clinical Trials: Datasets collected from controlled experiments that often include detailed patient information.

Features

- Characteristics used for prediction, such as:
- Symptoms (e.g., fever, cough)
- Medical history (e.g., previous illnesses, surgeries)
- Laboratory results (e.g., blood pressure, glucose levels)



Key Components of Medical Diagnosis with Supervised Learning

Labels

- The target variable that represents the diagnosis (e.g., “diabetes,” “cancer,” “no disease”).

Model Training

- Algorithms are trained on historical data where the outcome (diagnosis) is known.
- Common algorithms include:
 - Logistic Regression, Decision Trees, Random Forests, SVM, Artificial Neural Networks.

Diabetes Prediction Example


Problem

- Predicting whether a patient has diabetes based on factors like age, BMI, blood pressure, and glucose levels, ...

How it Works

- A **logistic regression model** is trained on a dataset of patients, learning to associate specific feature values with the presence of diabetes.

Prediction Model of Diabetes



Dataset selection and feature evaluation



- The Pima Indians Diabetes Dataset (**PIMA**) is a publicly available test dataset widely used for diabetes research and predictive modeling.
- It contains **768 observations** of females of Pima Indian heritage **aged 21 years or older**.
- The dataset includes **eight medical predictor variables** and **one target variable**.



Dataset selection and feature evaluation (predictor variables)

- **Pregnancies**: Number of times pregnant
- **Glucose**: Plasma glucose concentration over 2 hours in an oral glucose tolerance test
- **Blood Pressure**: Diastolic blood pressure (mm Hg)
- **Skin Thickness**: Triceps skinfold thickness (mm)
- **Insulin**: 2-hour serum insulin (mu U/ml)
- **BMI**: Body mass index
- **Diabetes** **Pedigree Function**: A function that scores the likelihood of diabetes based on family history
- **Age**: Age in years



Dataset selection and feature evaluation (target variable)



- The target variable is '**Outcome**' which indicates whether the patient had diabetes (1) or not (0).
- This training dataset is particularly useful for testing machine learning algorithms for **binary classification tasks**.



Dataset selection and feature evaluation (training dataset)



	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1




Correlation between all the features



- Analyze the correlation between all the features in the dataset before any data cleaning.
- This will help us understand the relationships between different features:





Correlation between all the features



- Inference: The heatmap above shows the correlation coefficients between each pair of features in the dataset. The correlation coefficient ranges from -1 to 1, where:
 - ✓ 1 indicates a perfect positive correlation,
 - ✓ -1 indicates a perfect negative correlation,
 - ✓ 0 indicates no correlation.



Correlation between all the features (Scaling the Data)



- Standard scaling helps normalize the data, ensuring that each feature contributes equally to the machine-learning model:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.639947	0.865108	-0.033518	0.670643	-0.181541	0.166619	0.468492	1.425995
1	-0.844885	-1.206162	-0.529859	-0.012301	-0.181541	-0.852200	-0.365061	-0.190672
2	1.233880	2.015813	-0.695306	-0.012301	-0.181541	-1.332500	0.604397	-0.105584
3	-0.844885	-1.074652	-0.529859	-0.695245	-0.540642	-0.633881	-0.920763	-1.041549
4	-1.141852	0.503458	-2.680669	0.670643	0.316566	1.549303	5.484909	-0.020496



Correlation between all the features (Target Column)



- The 'Outcome' column shows whether a patient has diabetes (1) or not (0). Understanding the target variable is essential to build an accurate predictive model.

Random Forest

Training Accuracy = 1.0

Test Accuracy = 0.77

Decision Tree

Training Accuracy = 1.0

Test Accuracy = 0.73

XgBoost Classifier

Training Accuracy = 1.0

Test Accuracy = 0.74


SVM

Training Accuracy = 1.0

Test Accuracy = 0.74



KMeans Clustering



Clustering

Definition: Clustering is the task of grouping similar items together based on their characteristics without prior knowledge of the groups.



Key Concepts in Clustering

Input Data

- Collect data that includes multiple features.
- This could be anything from customer data (age, income, purchase history) to image data (pixel values).

Feature Selection

- Identify the features that will be used for clustering.
- It's essential to choose relevant features that help distinguish the different groups.

Key Concepts in Clustering

Clustering Algorithms

- **K-Means Clustering:** Partitions data into a specified number of clusters (k), minimizing the variance within each cluster.
- **Hierarchical Clustering:** Creates a tree of clusters, allowing for different levels of granularity.
- **DBSCAN:** Groups together points that are closely packed while marking points in low-density regions as outliers.
- **Gaussian Mixture Models (GMM):** Assumes data points are generated from a mixture of several Gaussian distributions.

Key Concepts in Clustering



The output

- The final output is the cluster assignments for each data point, allowing you to see which items are grouped together.



Brain tumor segmentation using clustering



Is a common application in medical imaging that helps identify and delineate tumor regions from brain MRI scans.



Key Components of Brain tumor segmentation using clustering

Data Acquisition

- Obtain MRI scans of the brain, which can be in different modalities (e.g., T1, T2).
- The dataset may include labeled images where the tumor regions are marked, but the clustering process does not require labeled data.

Preprocessing

- **Image Enhancement:** Improve image quality using techniques like histogram equalization.
- **Noise Reduction:** Apply filters to reduce noise in the images.
- **Normalization:** Scale pixel intensities for uniformity across different scans.

Key Components of Brain tumor segmentation using clustering




Feature Extraction

- Convert the 2D MRI images into features suitable for clustering. Common features include:
 - Pixel Intensity
 - Texture Features
 - Spatial Information


Clustering Algorithm

- K-Means Clustering
- Fuzzy C-Means
- DBSCAN
- Hierarchical Clustering





- Proposed FPCM-based tumor segmentation on brain MR datasets (I-01, I-02, I-03, I-04, and I-05).
- The first column denotes original images,
- The second column denotes segmented tumors mask, and
- The third column denotes tumors delineated by the proposed method (yellow) and manual segmentation (red).



Prediction

Definition: Prediction involves using historical data to forecast future outcomes or trends.

This can apply to various fields, such as finance, healthcare, marketing, and more.



Key Concepts in Prediction

Data Collection

- Gather relevant historical data that pertains to the outcome you wish to predict. This data can come from various sources, including databases, APIs, or even user-generated content.

Data Preprocessing

- **Cleaning:** Remove missing or erroneous data points.
- **Normalization:** Scale the data for consistency,
- **Feature Engineering:** Create new features or transform existing ones to improve model performance.

Key Concepts in Prediction

Prediction Model

- Depending on the nature of the data and the type of prediction, you might choose from various models, such as:
 - Regression Models
 - Classification Models
 - Time Series Models
 - Machine Learning Algorithms

Model Training

- Split your data into training and testing sets.
- Train the model using the training data, adjusting parameters to minimize prediction errors.

Key Concepts in Prediction

The output (Prediction)

- Once the model is trained and validated, you can use it to make predictions on new, unseen data.

Predicting stock prices



Predicting stock prices involves analyzing past market data and trends to make informed guesses about future price movements



Predicting stock prices

Data Collection

- Gather historical stock price data, which typically includes:
- Open, High, Low, Close Prices: Daily or intraday trading data.
- Volume: The number of shares traded.
- Market Indicators: Data like moving averages or Bollinger Bands.
- Fundamental Data: Company earnings reports, economic indicators, and news articles.

Data

Preprocessing

- **Cleaning**: Remove any missing or erroneous entries in the dataset.
- **Normalization**: (e.g., neural networks).
- **Feature Engineering**: Create additional features that may help in predicting stock prices:
- **Technical Indicators**: Calculate moving averages, momentum indicators, etc.

Predicting stock prices

Prediction Model

- Time Series Analysis: A model like ARIMA is popular for analyzing temporal data.
- Machine Learning Models: Algorithms such as:
 - Linear Regression
 - Decision Trees or Random Forests:
 - Support Vector Machines (SVM): Effective for classification tasks in stock price movement (up/down).
 - Neural Networks: RNNs or LSTM are particularly suited for sequential data like stock prices.


Model Training


- Split the dataset into training and testing sets (e.g., 80/20 split).
- Train the model using the training set, adjusting parameters to minimize prediction errors.

Predicting stock prices

The output (Prediction)

- Use the trained model to predict future stock prices based on the latest available data.






05

Popular Machine Learning Frameworks



Challenges in Machine Learning

- **Data Quality**: Inaccurate or biased data can lead to poor model performance.
- **Overfitting**: When a model learns the noise in the training data rather than the underlying pattern.
- **Interpretability**: Difficulty in understanding complex models, especially deep learning models.



Thanks!

Do you have any questions?

